1- Classification Notebook:

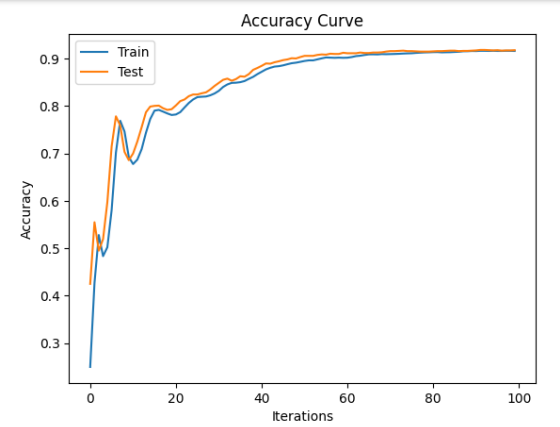

The model demonstrates strong accuracy performance without signs of overfitting. Both training and test accuracy curves converge around 90%, indicating the model achieves about 90% prediction accuracy consistently across data splits. This convergence highlights the model's ability to generalize well from the training data to unseen test data.

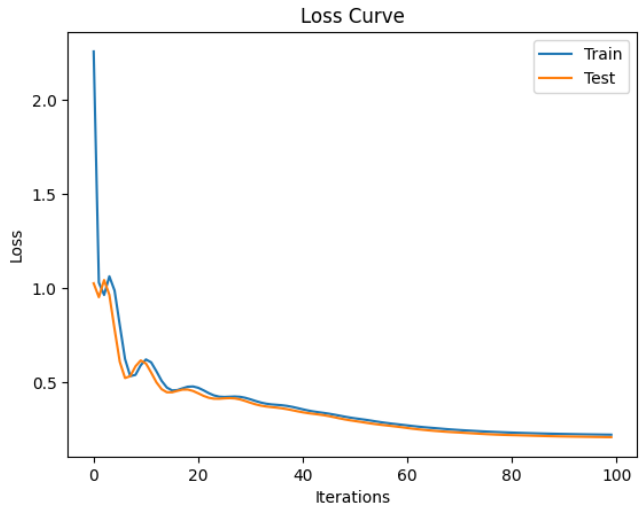

The loss curves for both training and testing data decrease steadily and closely follow one another. This alignment between train and test loss suggests stable model learning, where the predictive performance improves over iterations without deviating or diverging. The overall loss reduction illustrates successful optimization and effective learning from the training data.

2- Vision notebook:

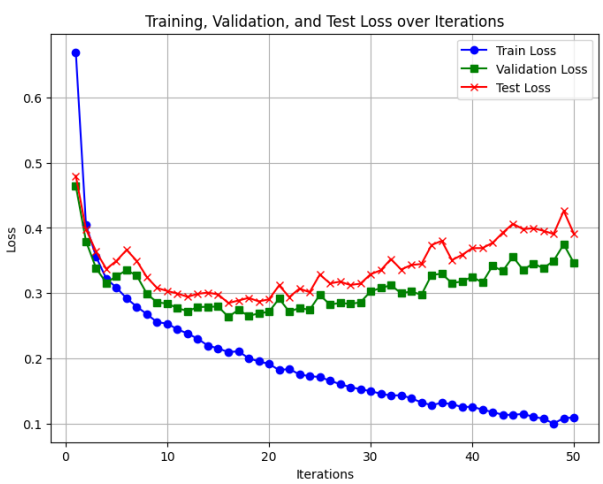

Training and test loss both steadily decrease and closely follow each other, indicating effective learning with no significant issues early on. However, looking at the training, validation, and test losses together reveals overfitting: training loss continues to decrease throughout, but validation and test losses start increasing after approximately 20 iterations, suggesting the model begins to memorize training data and generalizes less well over time.

The training and test accuracy curves show steady improvement, both reaching around 90% by the later iterations, indicating good learning and generalization. However, when including validation results, the training accuracy continues to rise towards 99%, while validation and test accuracies plateau around 91-92%, demonstrating slight overfitting as the model fits the training data better than unseen data.